

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score, mean_absolute_error, mean_squared_error
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestRegressor
import joblib
import matplotlib.pyplot as plt
```

```
In [2]: student_df = pd.read_csv("DataSets/student_scores.csv")
student_df
```

Out[2]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25
10	7.7	85
11	5.9	62
12	4.5	41
13	3.3	42
14	1.1	17
15	8.9	95
16	2.5	30
17	1.9	24
18	6.1	67
19	7.4	69
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86


Some Visualization

```
In [3]: scores = student_df["Scores"]
        hours = student_df["Hours"]

        # Create the bar plot
        plt.bar( hours,scores)


        # Set labels and title
        plt.xlabel('Scores')
        plt.ylabel('Hours')
        plt.title('Scores vs Hours')

        # Display the plot
        plt.show()
```



```
In [4]: scores = student_df["Scores"]  
hours = student_df["Hours"]  
  
# Create the bar plot  
plt.scatter(hours,scores);  
plt.xlabel("Hours")  
plt.ylabel("Scores")  
plt.title("Hours vs Scores")
```

Out[4]: Text(0.5, 1.0, 'Hours vs Scores')



In []:

Model Training

```
In [5]: x = np.array(student_df["Hours"]).reshape(-1,1)
y = student_df["Scores"]
```

```
In [6]: # Splitting the Data
x_train ,x_test ,y_train ,y_test = train_test_split(x,y,test_size=0.3, random_state=41)
```

```
In [7]: model = LinearRegression()
model.fit(x_train ,y_train)
```

```
Out[7]:
```

▼ LinearRegression

LinearRegression()

```
In [8]: model.score(x_test,y_test)
```

```
Out[8]: 0.9621346134566173
```

```
In [9]: joblib.dump(model,"TrainedModels/StudentScore.pkl")
```

```
Out[9]: ['TrainedModels/StudentScore.pkl']
```

Load the Model to make Prediction

```
In [10]: mod = joblib.load("TrainedModels/StudentScore.pkl")
```

```
In [11]: y_pred = mod.predict(x_test)

pd.DataFrame({"y_test": y_test, "y_pred": y_pred})
```

Out[11]:

	y_test	y_pred
5	20	14.990287
19	69	74.884076
14	17	10.929691
10	85	77.929523
7	60	55.596245
8	81	84.020416
17	24	19.050883
11	62	59.656841

```
In [12]: mse = mean_squared_error(y_test, y_pred)
print(f"mean squared error is {mse}")
mae = mean_absolute_error(y_test, y_pred)
print(f"mean absolute error is {mae}")
```

```
mean squared error is 25.63250010321114
mean absolute error is 4.843877805244228
```

In []:

In []:

Task 2 Petrol Consumption


```
In [13]: Petrol_df = pd.read_csv("DataSets/petrol_consumption.csv")
Petrol_df.head()
```

Out[13]:

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)	Petrol_Consumption
0	9.0	3571	1976	0.525	541
1	9.0	4092	1250	0.572	524
2	9.0	3865	1586	0.580	561
3	7.5	4870	2351	0.529	414
4	8.0	4399	431	0.544	410


Visualization

```
In [14]: target = Petrol_df['Petrol_Consumption']
Average_income = Petrol_df['Average_income']
Paved_Highways = Petrol_df['Paved_Highways']
Petrol_tax = Petrol_df['Petrol_tax']
Population_Driver_licence = Petrol_df['Population_Driver_licence(%)']

plt.subplot(2, 1, 1)
plt.scatter(target, Average_income, label='Average_income')
plt.scatter(target, Paved_Highways, label='Paved_Highways')
plt.xlabel('Target')
plt.ylabel('Feature')
plt.title('Scatter Plot')
plt.legend()

plt.subplot(2, 1, 2)
plt.scatter(target, Petrol_tax, label='Petrol_tax')
plt.scatter(target, Population_Driver_licence, label='Population_Driver_licence')
plt.xlabel('Target')
plt.ylabel('Feature')
plt.legend()

plt.tight_layout() # Adjusts the spacing between subplots
plt.show()
```



In []:

```

In [24]: #Use Cols Transformer for MinMax Scalling
minmax =MinMaxScaler()

transform_cols=["Petrol_tax","Average_income","Paved_Highways","Population_Driver_licence(%)"]

transformer =ColumnTransformer([("minmax",minmax,transform_cols)])

data = transformer.fit_transform(Petrol_df)

data = pd.DataFrame(data,columns=["Petrol_tax","Average_income","Paved_Highways","Population_Driver_licence(%)"]
data.head()

```

Out[24]:

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)
0	0.8	0.222905	0.089044	0.271062
1	0.8	0.451514	0.047202	0.443223
2	0.8	0.351909	0.066567	0.472527
3	0.5	0.792892	0.110656	0.285714
4	0.6	0.586222	0.000000	0.340659

```

In [16]: x = data
y = Petrol_df["Petrol_Consumption"]

```

```

In [17]: x_train ,x_test ,y_train ,y_test = train_test_split(x,y,test_size=0.3 ,random_state=52)

```

```

In [18]: model1 = LinearRegression()
model1.fit(x_train ,y_train)

```

Out[18]:

LinearRegression

LinearRegression()

```
In [19]: model1.score(x_test,y_test)
```

```
Out[19]: 0.7812556015921516
```

```
In [20]: joblib.dump(model1,"TrainedModels/PetrolConsumption.pkl")
```

```
Out[20]: ['TrainedModels/PetrolConsumption.pkl']
```

```
In [21]: mod1 = joblib.load("TrainedModels/PetrolConsumption.pkl")
```

```
In [22]: y_pred = mod1.predict(x_test)  
  
pd.DataFrame({"y_test": y_test, "y_pred": y_pred})
```

```
Out[22]:
```

	y_test	y_pred
8	464	498.887986
3	414	488.882860
25	566	549.211935
44	782	675.788092
18	865	745.446171
26	577	597.792995
6	344	337.080999
9	498	547.265145
34	487	520.150864
17	714	594.451750
1	524	561.120649
30	571	570.485001
0	541	540.348735
38	648	714.693951
21	540	559.453224

```
In [23]: mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")

mae = mean_absolute_error(y_test, y_pred)
print(f"Mean Absolute Error: {mae}")
```

Mean Squared Error: 3812.248209532195
Mean Absolute Error: 47.09566602370284

In []:

In []: