

```
In [1]: # Import Libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.compose import ColumnTransformer
import joblib
```

Task # 1

In dataset "Fish.csv", Take Species, Length, Width as Input and predict its height using Linear Regression. You can validate your model through Descent Gradient approach to check results.

```
In [2]: fish_df = pd.read_csv("DataSets/Fish.csv")
fish_df.head()
```

Out[2]:

| | Species | Weight | Length1 | Length2 | Length3 | Height | Width |
|---|---------|--------|---------|---------|---------|---------|--------|
| 0 | Bream | 242.0 | 23.2 | 25.4 | 30.0 | 11.5200 | 4.0200 |
| 1 | Bream | 290.0 | 24.0 | 26.3 | 31.2 | 12.4800 | 4.3056 |
| 2 | Bream | 340.0 | 23.9 | 26.5 | 31.1 | 12.3778 | 4.6961 |
| 3 | Bream | 363.0 | 26.3 | 29.0 | 33.5 | 12.7300 | 4.4555 |
| 4 | Bream | 430.0 | 26.5 | 29.0 | 34.0 | 12.4440 | 5.1340 |

cLEAN AND PROCESS THE DATA

```
In [3]: fish_df.isna().sum()    # so no need to clean
```

```
Out[3]: Species      0  
        Weight      0  
        Length1     0  
        Length2     0  
        Length3     0  
        Height      0  
        Width       0  
        dtype: int64
```

```
In [4]: fish_df.drop_duplicates().shape
```

```
Out[4]: (159, 7)
```

```
In [ ]:
```



```
In [5]: import matplotlib.pyplot as plt
import pandas as pd

fish_df = pd.read_csv("DataSets/Fish.csv")
height = fish_df['Height']
weight = fish_df['Weight']
length1 = fish_df['Length1']
length2 = fish_df['Length2']
length3 = fish_df['Length3']
width = fish_df['Width']

# Set the figure size
fig, axes = plt.subplots(2, 3, figsize=(12, 8))

# Plot Height vs. Weight
axes[0, 0].scatter(weight, height)
axes[0, 0].set_xlabel('Weight')
axes[0, 0].set_ylabel('Height')
axes[0, 0].set_title('Height vs. Weight')

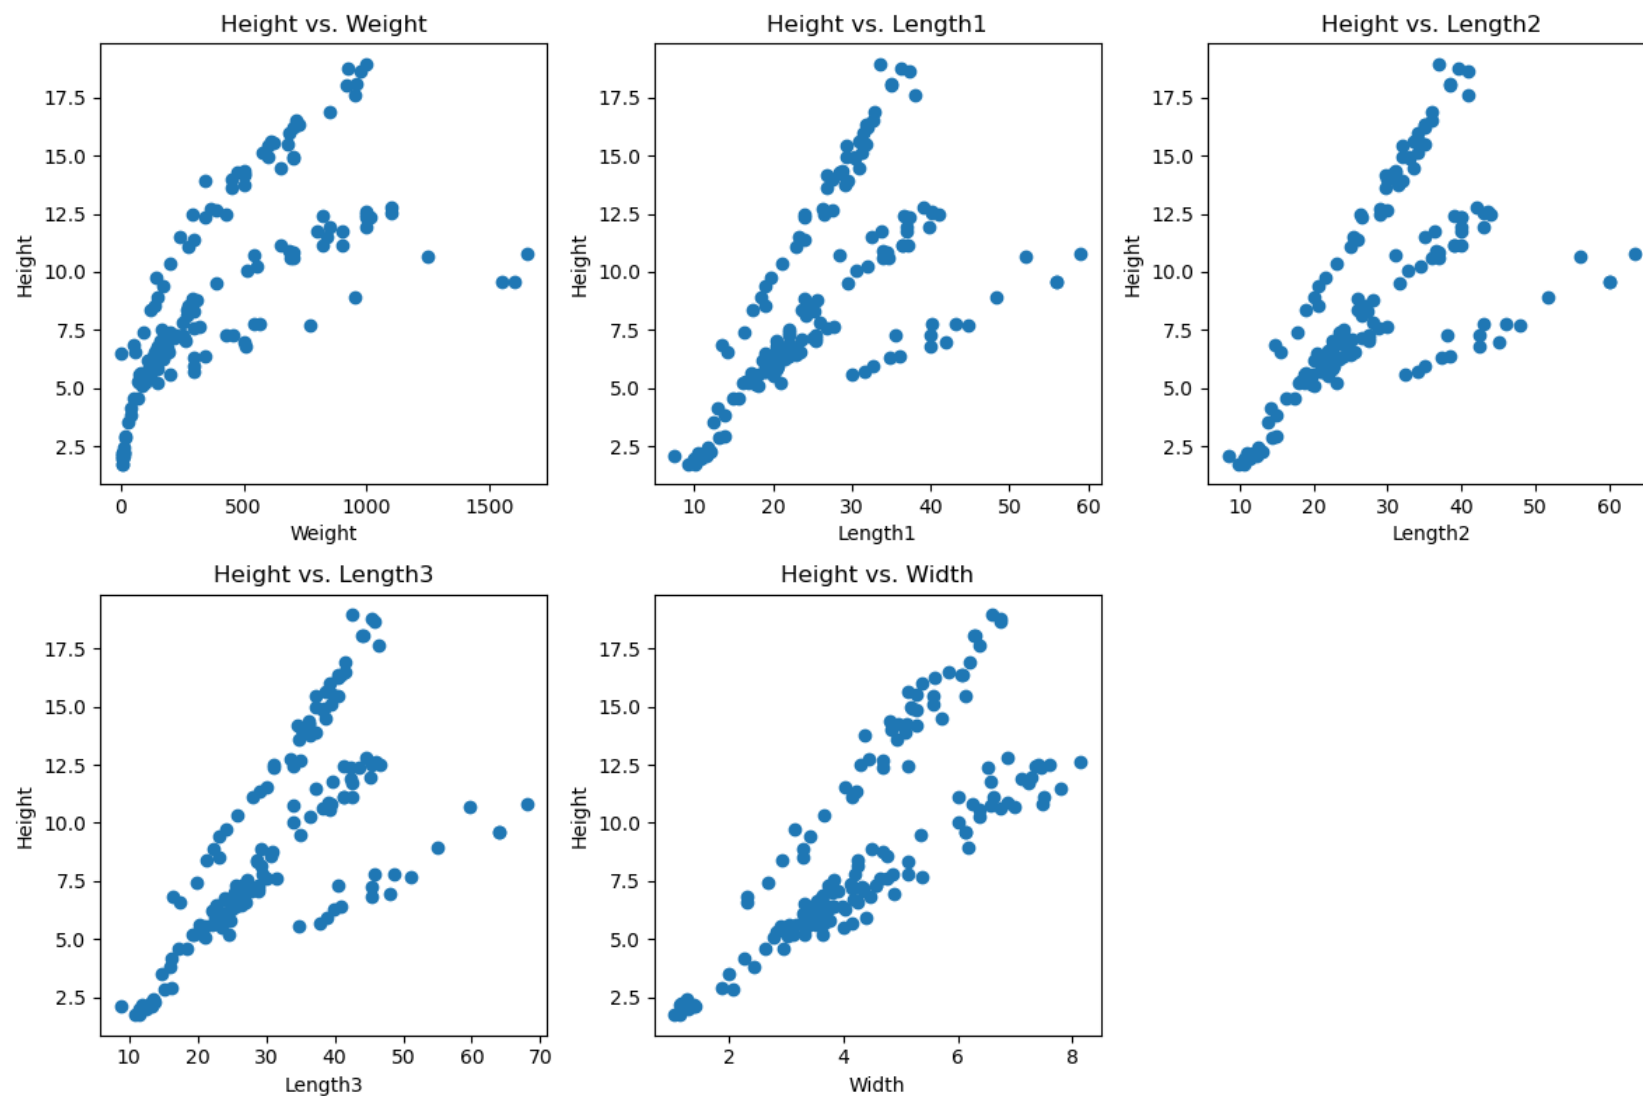
# Plot Height vs. Length1
axes[0, 1].scatter(length1, height)
axes[0, 1].set_xlabel('Length1')
axes[0, 1].set_ylabel('Height')
axes[0, 1].set_title('Height vs. Length1')

# Plot Height vs. Length2
axes[0, 2].scatter(length2, height)
axes[0, 2].set_xlabel('Length2')
axes[0, 2].set_ylabel('Height')
axes[0, 2].set_title('Height vs. Length2')

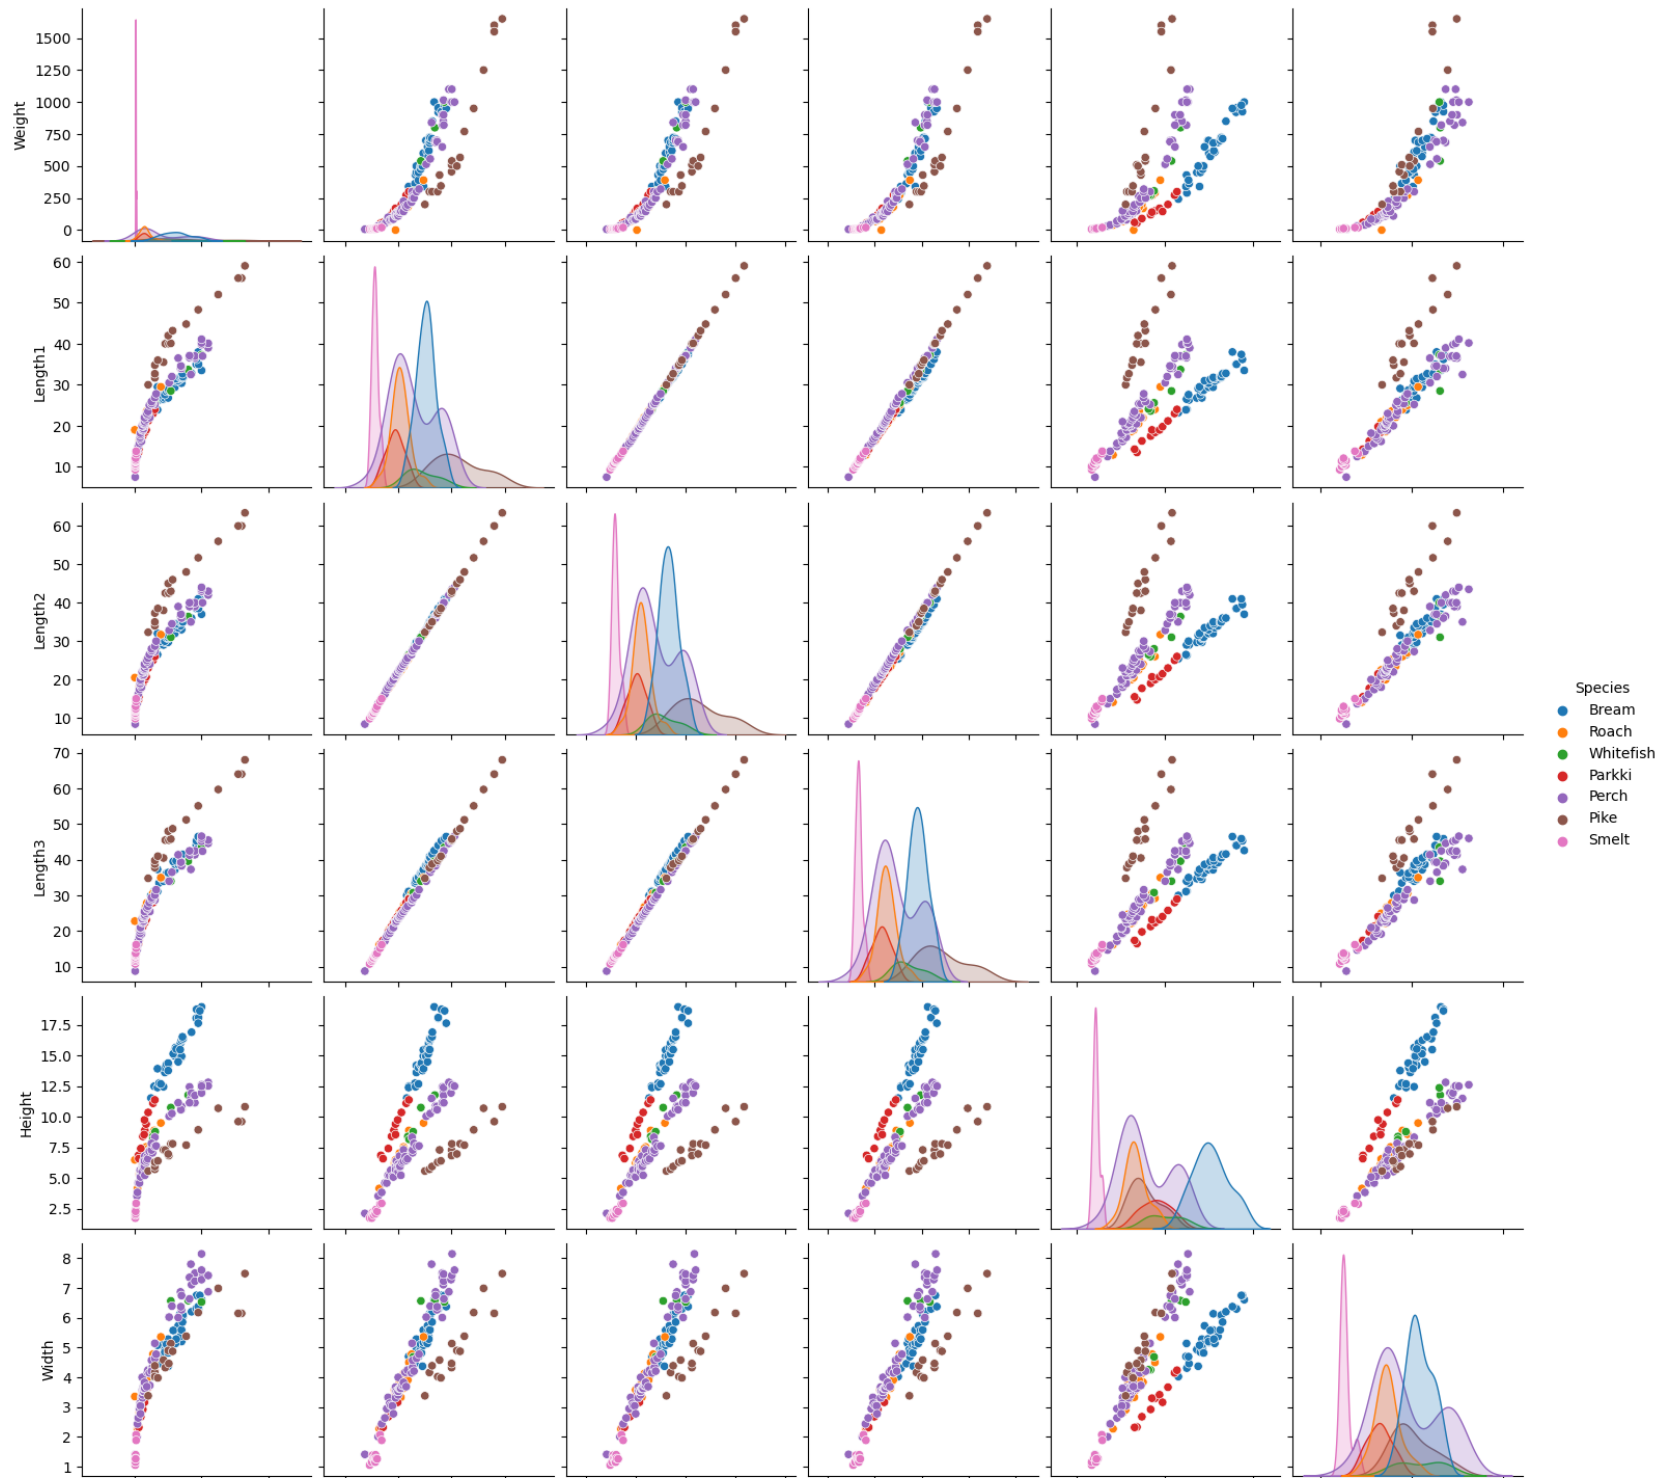
# Plot Height vs. Length3
axes[1, 0].scatter(length3, height)
axes[1, 0].set_xlabel('Length3')
axes[1, 0].set_ylabel('Height')
axes[1, 0].set_title('Height vs. Length3')

# Plot Height vs. Width
axes[1, 1].scatter(width, height)
axes[1, 1].set_xlabel('Width')
axes[1, 1].set_ylabel('Height')
axes[1, 1].set_title('Height vs. Width')
```

```
# Remove empty subplot  
axes[1, 2].remove()  
  
plt.tight_layout()  
plt.show()
```



```
In [6]: import seaborn as sns  
sns.pairplot(data=fish_df, hue="Species");
```



In []:

First Apply basic model to evaluate the Dataset

```
In [7]: x = fish_df.drop("Height",axis=1)
y = fish_df["Height"]
x["Species"].value_counts()
```

```
Out[7]: Perch      56
Bream      35
Roach      20
Pike       17
Smelt      14
Parkki     11
Whitefish   6
Name: Species, dtype: int64
```

```
In [8]: category = ["Species"]
encode = OneHotEncoder()

transf = ColumnTransformer([("encode",encode,category)],remainder="passthrough")
transformed_x = transf.fit_transform(x)
transformed_x
```

```
Out[8]: array([[ 1.    ,  0.    ,  0.    , ..., 25.4   , 30.    ,  4.02   ],
 [ 1.    ,  0.    ,  0.    , ..., 26.3   , 31.2   ,  4.3056 ],
 [ 1.    ,  0.    ,  0.    , ..., 26.5   , 31.1   ,  4.6961 ],
 ...,
 [ 0.    ,  0.    ,  0.    , ..., 13.    , 13.8   ,  1.2558 ],
 [ 0.    ,  0.    ,  0.    , ..., 14.3   , 15.2   ,  2.0672 ],
 [ 0.    ,  0.    ,  0.    , ..., 15.    , 16.2   ,  1.8792 ]])
```

```
In [9]: x_train ,x_test ,y_train ,y_test = train_test_split(transformed_x,y,test_size=0.2,random_state=42)
```

```
In [10]: model = LinearRegression()  
         model.fit(x_train,y_train)
```

```
Out[10]: 

▼ LinearRegression



LinearRegression()


```

```
In [11]: model.score(x_test,y_test)
```

```
Out[11]: 0.980397551332076
```

```
In [12]: y_pred = model.predict(x_test)
         pd.DataFrame({"y_test":y_test , "y_pred":y_pred})
```

Out[12]:

| | y_test | y_pred |
|------------|---------|-----------|
| 78 | 5.1992 | 4.999679 |
| 155 | 2.4300 | 2.130787 |
| 128 | 5.5680 | 4.321885 |
| 55 | 8.3804 | 7.996653 |
| 94 | 5.2185 | 6.162787 |
| 29 | 18.9570 | 16.658972 |
| 147 | 1.7284 | 1.769991 |
| 51 | 7.0866 | 7.286182 |
| 98 | 6.7334 | 6.927047 |
| 141 | 10.6863 | 10.812777 |
| 19 | 14.4738 | 15.310182 |
| 60 | 12.3540 | 11.922111 |
| 15 | 15.4380 | 15.008784 |
| 65 | 8.8928 | 9.002424 |
| 24 | 16.2405 | 15.410072 |
| 30 | 18.0369 | 16.536581 |
| 126 | 12.6040 | 13.274712 |
| 101 | 7.1680 | 7.052905 |
| 96 | 7.2930 | 6.351509 |
| 16 | 14.8604 | 14.789092 |
| 151 | 2.2139 | 2.109468 |
| 18 | 15.6330 | 14.673377 |
| 12 | 13.7592 | 13.614805 |
| 9 | 14.2266 | 14.270440 |
| 31 | 18.0840 | 16.488222 |
| 125 | 12.5125 | 12.380040 |

| | y_test | y_pred |
|------------|---------------|---------------|
| 95 | 6.2750 | 6.317563 |
| 56 | 8.1454 | 8.073877 |
| 145 | 1.7388 | 1.572797 |
| 152 | 2.2139 | 1.983054 |
| 135 | 6.8250 | 6.625377 |
| 76 | 4.5880 | 4.732026 |

Gradient Decsent Method

```
In [13]: import numpy as np
cost = float('inf')
def gradient_Descent(X,Y):
    m_curr = c_curr =0
    iteration = 10000
    LearningRate = 0.00001
    n = len(X)
    cost = float('inf')
    i=0
    for i in range(iteration):

        cost_prev_val = cost
        y_predicted = (m_curr * X) + c_curr
        cost = 1/n * np.sum([val**2 for val in (Y - y_predicted)])
        mp = -(2/n) * np.sum(X*(Y-y_predicted))

        cp = -(2/n) * np.sum(Y-y_predicted)
        m_curr = m_curr - LearningRate * mp
        c_curr = c_curr - LearningRate * cp

        print("m_curr =" , m_curr)
        print("c_curr =" , c_curr)
        print("cost =" , cost)
        print("\nitteration =" , i)
```

```
In [19]: X =np.array([4 ,8,9,10,12,14,15])
Y =np.array([12 ,15,17,20,25,30,35])
```

```
In [20]: gradient_Descent(X,Y)

ititeration = 9977
m_curr = 2.122044512011252
c_curr = 0.18312784064554202
cost = 4.609652226846168

ititeration = 9978
m_curr = 2.1220445292356866
c_curr = 0.1831276432176868
cost = 4.609652222918709

ititeration = 9979
m_curr = 2.1220445464600854
c_curr = 0.18312744579023682
cost = 4.609652218991269

ititeration = 9980
m_curr = 2.1220445636844487
c_curr = 0.18312724836319208
cost = 4.609652215063812
```

```
In [ ]:
```

Task 2

In dataset "Salary.csv", Take no. of years as Input and Predict Employee Salary. Since its 2D dataset, check and validate model using Linear Regression by Gradient descent approach

LinearRegression model

In [126]:

```
# Import Libraries
import numpy as np
import pandas as pd
from sklearn.metrics import accuracy_score, mean_absolute_error, mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.compose import ColumnTransformer
import joblib
```

In [127]: `Salary_df = pd.read_csv("DataSets/Salary_dataset.csv")`
`Salary_df.head()`

Out[127]:

| | Unnamed: 0 | YearsExperience | Salary |
|---|------------|-----------------|---------|
| 0 | 0 | 1.2 | 39344.0 |
| 1 | 1 | 1.4 | 46206.0 |
| 2 | 2 | 1.6 | 37732.0 |
| 3 | 3 | 2.1 | 43526.0 |
| 4 | 4 | 2.3 | 39892.0 |

In [128]: `x = np.array(Salary_df["YearsExperience"]).reshape(-1, 1)`
`y = np.array(Salary_df["Salary"])`

In [129]: `x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=42)`

In [130]: `model = LinearRegression()`
`model.fit(x_train, y_train)`

Out[130]:

```
▼ LinearRegression
LinearRegression()
```



```
In [131]: model.score(x_test,y_test)
```

```
Out[131]: 0.9414466227178215
```

```
In [132]: y_pred = model.predict(x_test)

pd.DataFrame({"y_test":y_test , "y_pred_using_model":y_pred})
```

```
Out[132]:
```

| | y_test | y_pred_using_model |
|---|----------|--------------------|
| 0 | 112636.0 | 115574.622884 |
| 1 | 67939.0 | 71680.938782 |
| 2 | 113813.0 | 102499.908470 |
| 3 | 83089.0 | 75416.571471 |
| 4 | 64446.0 | 55804.499851 |
| 5 | 57190.0 | 60474.040713 |
| 6 | 122392.0 | 122111.980090 |
| 7 | 109432.0 | 107169.449332 |
| 8 | 56958.0 | 63275.765230 |

```
In [133]: mse = mean_squared_error(x_train,y_train)
print(f"Mean Squared Error: {mse}")

mae = mean_absolute_error(y_test, y_pred)
print(f"Mean Absolute Error: {mae}")
```

```
Mean Squared Error: 5728211930.733333
Mean Absolute Error: 5161.328710400178
```

Gradient Descent

```
In [136]: import numpy as np
cost = float('inf')
def gradient_Descent(X,Y):
    m_curr = c_curr =0
    iteration = 1000
    LearningRate = 0.01
    n = len(X)

    i=0
    for i in range(iteration):

        y_predicted = (m_curr * X) + c_curr
        cost = 1/n * np.sum([val**2 for val in (Y - y_predicted)])
        mp = -(2/n) * np.sum(X*(Y-y_predicted))

        cp = -(2/n) * np.sum(Y-y_predicted)
        m_curr = m_curr - LearningRate * mp
        c_curr = c_curr - LearningRate * cp

        print("m_curr =" , m_curr)
        print("c_curr =" , c_curr)
        print("cost =" , cost)
        print("\niteration =" , i)
```

```
In [137]: gradient_Descent(x_train.reshape(1,-1)[0],y_train.reshape(1,-1)[0])
```

```
ititeration = 990  
m_curr = 9376.779222394893  
c_curr = 24744.284819751472  
cost = 28936001.116909843
```

```
ititeration = 991  
m_curr = 9376.606487117215  
c_curr = 24745.390240080145  
cost = 28935875.072082177
```

```
ititeration = 992  
m_curr = 9376.434543337044  
c_curr = 24746.49059521631  
cost = 28935750.17971855
```

```
ititeration = 993  
m_curr = 9376.263387427627  
c_curr = 24747.585908369394  
cost = 28935626.429281704
```

```
In [138]: m = 9375.586570542466  
          c = 24752.98764930146
```

```
In [139]: y = m*x + c
```

```
In [140]: x_values = x_test.reshape(1,-1)[0]

y_values = []

for x in x_values:
    y = m * x + c
    y_values.append(y)

print("Resulting y values:", y_values)
```

Resulting y values: [115696.17738356336, 71630.92050201379, 102570.35618480391, 75381.15513023078, 55692.423332091596, 60380.21661736283, 122259.0879829431, 107258.14947007515, 63192.89258852557]

```
In [141]: pd.DataFrame({"y_test":y_test , "y_pred_Gradient":y_values})
```

Out[141]:

| | y_test | y_pred_Gradient |
|---|----------|-----------------|
| 0 | 112636.0 | 115696.177384 |
| 1 | 67939.0 | 71630.920502 |
| 2 | 113813.0 | 102570.356185 |
| 3 | 83089.0 | 75381.155130 |
| 4 | 64446.0 | 55692.423332 |
| 5 | 57190.0 | 60380.216617 |
| 6 | 122392.0 | 122259.087983 |
| 7 | 109432.0 | 107258.149470 |
| 8 | 56958.0 | 63192.892589 |

In []:

In []:

