

## Task 1

```
In [221]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder ,OneHotEncoder ,StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score ,confusion_matrix ,classification_report
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
In [222]: df = pd.read_csv("../DataSets/heart.csv")
df.head(5)
```

Out[222]:

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisea
0	40	M	ATA	140	289	0	Normal	172	N	0.0	Up	
1	49	F	NAP	160	180	0	Normal	156	N	1.0	Flat	
2	37	M	ATA	130	283	0	ST	98	N	0.0	Up	
3	48	F	ASY	138	214	0	Normal	108	Y	1.5	Flat	
4	54	M	NAP	150	195	0	Normal	122	N	0.0	Up	

```
In [223]: ##split the data into x and y

x = df.drop("HeartDisease",axis = 1)
y = df["HeartDisease"]
x.head(1)
```

Out[223]:

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope
0	40	M	ATA	140	289	0	Normal	172	N	0.0	Up

```
In [224]: ## standard deviation
std_devs = df.std()
print(std_devs)
```

```
Age          9.432617
RestingBP    18.514154
Cholesterol  109.384145
FastingBS     0.423046
MaxHR        25.460334
Oldpeak      1.066570
HeartDisease 0.497414
dtype: float64
```

C:\Users\Barcha\AppData\Local\Temp\ipykernel\_13644\1787614816.py:2: FutureWarning: The default value of numeric\_only in DataFrame.std is deprecated. In a future version, it will default to False. In addition, specifying 'numeric\_only=None' is deprecated. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
std_devs = df.std()
```

```
In [183]: ## Will apply all the math function in it  
df.describe()
```

Out[183]:

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	HeartDisease
<b>count</b>	918.000000	918.000000	918.000000	918.000000	918.000000	918.000000	918.000000
<b>mean</b>	53.510893	132.396514	198.799564	0.233115	136.809368	0.887364	0.553377
<b>std</b>	9.432617	18.514154	109.384145	0.423046	25.460334	1.066570	0.497414
<b>min</b>	28.000000	0.000000	0.000000	0.000000	60.000000	-2.600000	0.000000
<b>25%</b>	47.000000	120.000000	173.250000	0.000000	120.000000	0.000000	0.000000
<b>50%</b>	54.000000	130.000000	223.000000	0.000000	138.000000	0.600000	1.000000
<b>75%</b>	60.000000	140.000000	267.000000	0.000000	156.000000	1.500000	1.000000
<b>max</b>	77.000000	200.000000	603.000000	1.000000	202.000000	6.200000	1.000000


```
In [184]: heCorr = df.corr()  
plt.figure(figsize = (14,14))  
sns.heatmap(heCorr , annot = True)
```


C:\Users\Barcha\AppData\Local\Temp\ipykernel\_13644\2068051862.py:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
heCorr = df.corr()
```

```
Out[184]: <Axes: >
```







```
In [185]: df['FastingBS'].value_counts()
```

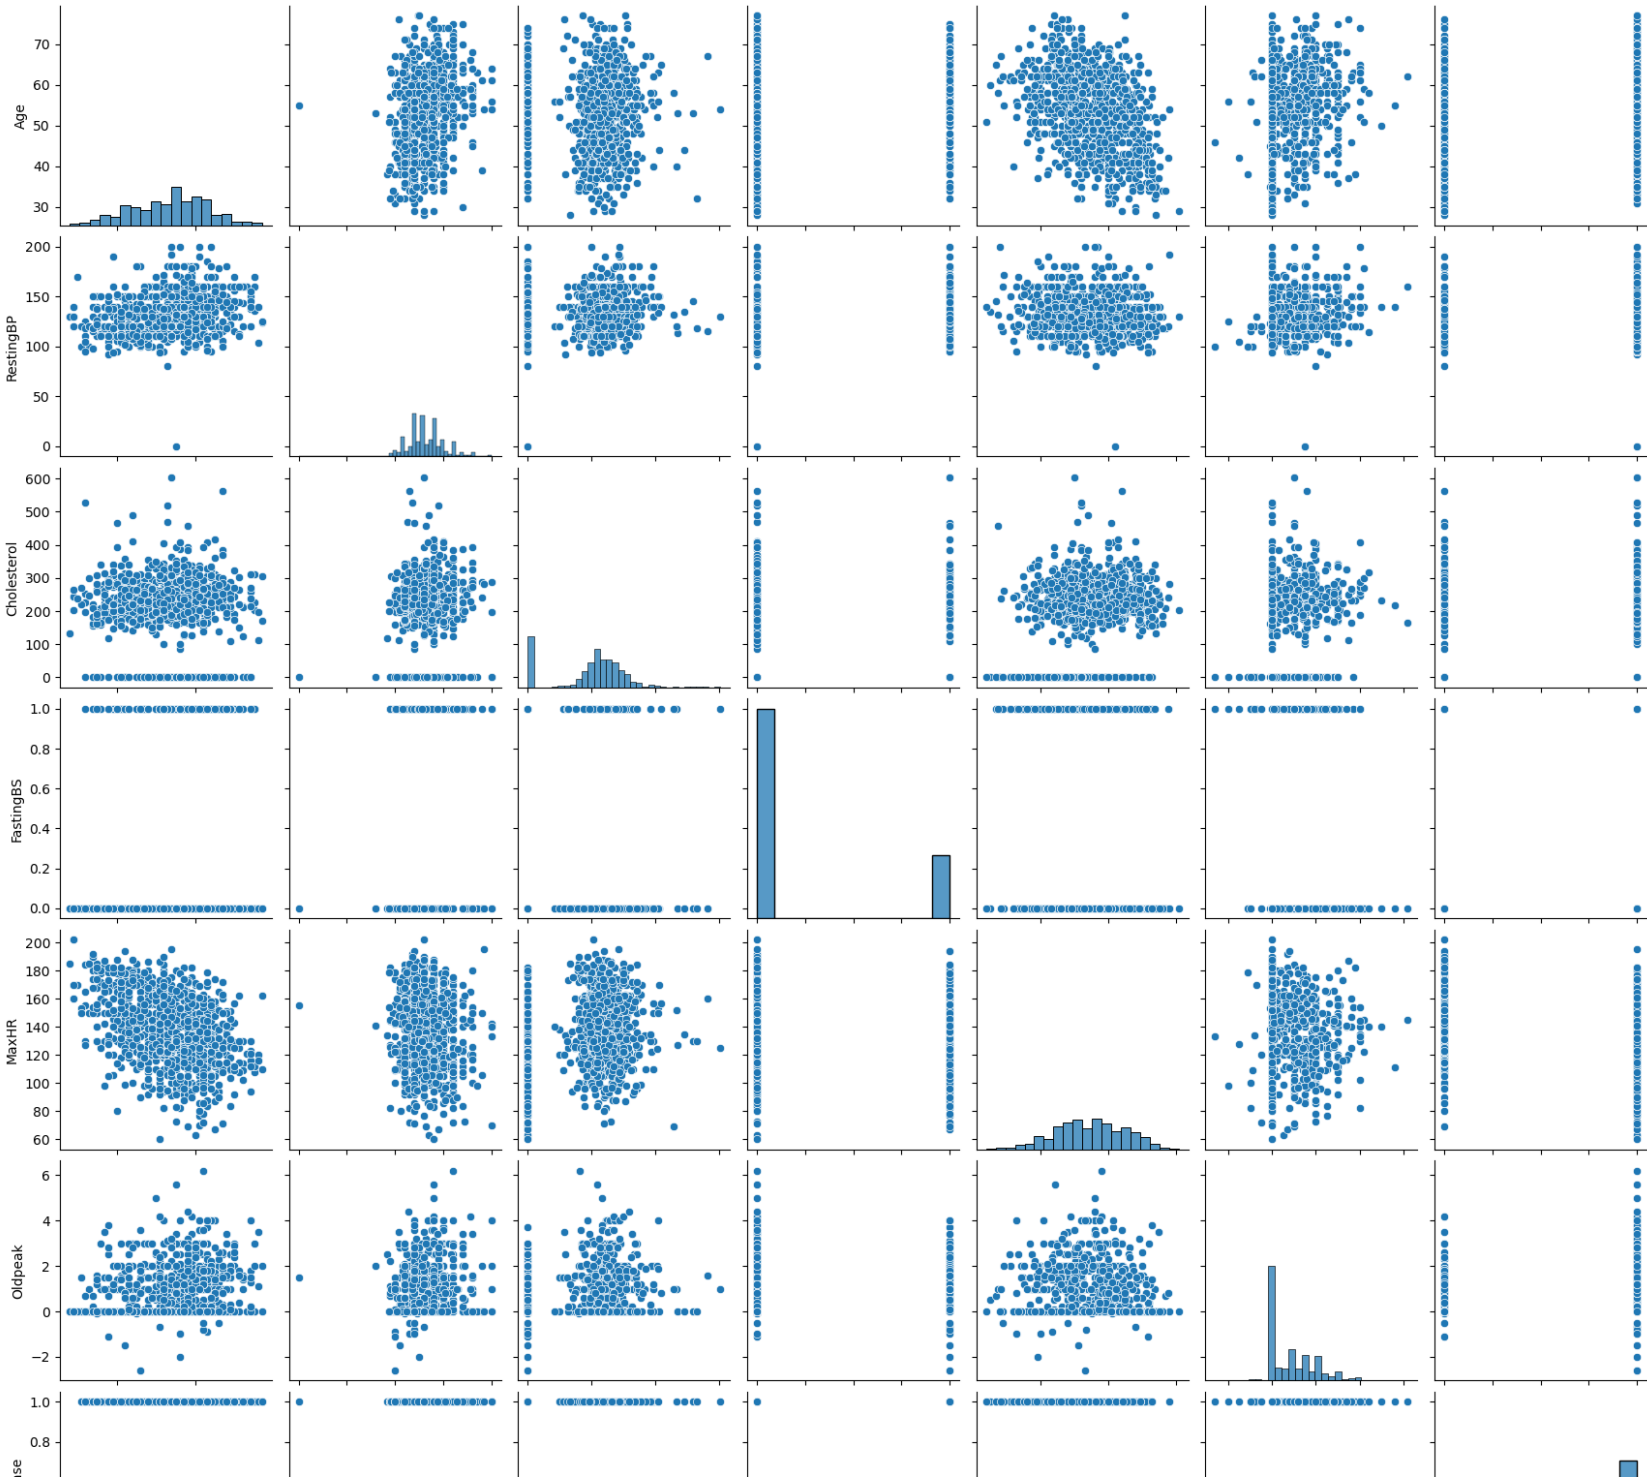
```
Out[185]: 0    704  
          1    214  
          Name: FastingBS, dtype: int64
```


```
In [225]: sns.pairplot(df)
```

```
Out[225]: <seaborn.axisgrid.PairGrid at 0x18eba94ff40>
```










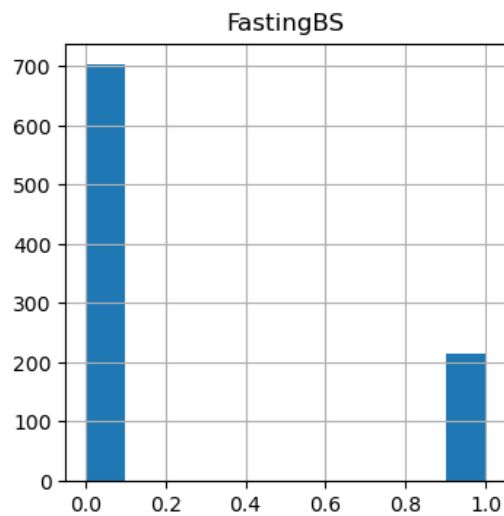

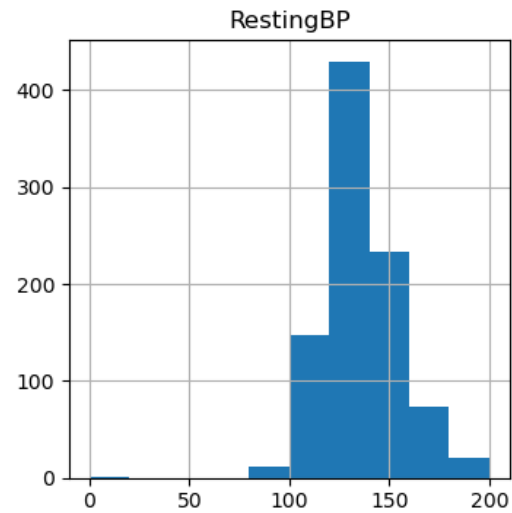




In [229]:

```
df.hist(figsize = [14,14])
```

```
Out[229]: array([[<Axes: title={'center': 'Age'}>,
  <Axes: title={'center': 'RestingBP'}>,
  <Axes: title={'center': 'Cholesterol'}>],
 [ <Axes: title={'center': 'FastingBS'}>,
  <Axes: title={'center': 'MaxHR'}>,
  <Axes: title={'center': 'Oldpeak'}>],
 [ <Axes: title={'center': 'HeartDisease'}>, <Axes: >, <Axes: >]],
 dtype=object)
```







In [186]: *## Data preprocessing*

```
encoder = OneHotEncoder()
scaling = StandardScaler()

categorical = ["Sex", "ChestPainType", "RestingECG", "ExerciseAngina", "ST_Slope"]
numerical = ["Age", "RestingBP", "Cholesterol", "MaxHR", "FastingBS", "Oldpeak"]

transform = ColumnTransformer([("categ", encoder, categorical),
                               ("scaling", scaling, numerical)], remainder = "passthrough")

x_trans = transform.fit_transform(x)
x_trans
```

```
Out[186]: array([[ 0.          ,  1.          ,  0.          , ...,  1.38292822,
                  -0.55134134, -0.83243239],
                 [ 1.          ,  0.          ,  0.          , ...,  0.75415714,
                  -0.55134134,  0.10566353],
                 [ 0.          ,  1.          ,  0.          , ..., -1.52513802,
                  -0.55134134, -0.83243239],
                 ...,
                 [ 0.          ,  1.          ,  1.          , ..., -0.85706875,
                  -0.55134134,  0.29328271],
                 [ 1.          ,  0.          ,  0.          , ...,  1.4615246 ,
                  -0.55134134, -0.83243239],
                 [ 0.          ,  1.          ,  0.          , ...,  1.42222641,
                  -0.55134134, -0.83243239]])
```

In [187]: *## Split the data into train test*

```
np.random.seed(42)
x_train ,x_test ,y_train ,y_test = train_test_split(x_trans ,y ,test_size = 0.3 ,random_state = 43)
```

# SVC Model

```
In [188]: ## model selection  
  
svc = SVC(C=0.5 ,kernel = "rbf")  
svc.fit(x_train ,y_train)  
svc.score(x_test ,y_test)
```

```
Out[188]: 0.855072463768116
```



```
In [189]: ## Hyperparameter tuning

from sklearn.model_selection import GridSearchCV

# defining parameter range
param_grid = {'C': [0.1, 1, 10, 500],
              'gamma': [1, 0.1, 0.001, 0.0001],
              'kernel': ['rbf', "poly", "linear"]}

grid = GridSearchCV(svc, param_grid, verbose = 2 ,cv = 3)

# fitting the model for grid search
svcHyper = grid.fit(x_train, y_train)
```

```
Fitting 3 folds for each of 48 candidates, totalling 144 fits
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=1, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=0.1, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=0.001, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=0.0001, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.0001, kernel=rbf; total time= 0.0s
[CV] END .....C=0.1, gamma=0.0001, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=0.0001, kernel=poly; total time= 0.0s
[CV] END .....C=0.1, gamma=0.0001, kernel=linear; total time= 0.0s
[CV] END .....C=0.1, gamma=0.0001, kernel=linear; total time= 0.0s
```

```
In [190]: svcHyper.best_score_
```

```
Out[190]: 0.8878504672897196
```

```
In [191]: svcHyper.best_params_
```

```
Out[191]: {'C': 1, 'gamma': 0.1, 'kernel': 'rbf'}
```

```
In [192]: ## y_pred  
y_pred = svc.predict(x_test)  
print(accuracy_score(y_test ,y_pred))
```


```
0.855072463768116
```

```
In [193]: sm = confusion_matrix(y_test ,y_pred)  
sm
```

```
Out[193]: array([[ 94,  22],  
                [ 18, 142]], dtype=int64)
```

```
In [194]: sns.heatmap(sm ,annot = True)
```

```
Out[194]: <Axes: >
```



## DecisionTreeClassifier

```
In [195]: ## Use decion tree model  
dec = DecisionTreeClassifier()  
dec.fit(x_train ,y_train)  
dec.score(x_test ,y_test)
```

```
Out[195]: 0.7463768115942029
```

```
In [196]: ## y_pred  
y_pred = dec.predict(x_test)  
print(accuracy_score(y_test ,y_pred))
```


0.7463768115942029

```
In [197]: s = confusion_matrix(y_test ,y_pred)  
s
```

```
Out[197]: array([[ 84,  32],  
               [ 38, 122]], dtype=int64)
```

```
In [198]: sns.heatmap(s ,annot = True)
```

```
Out[198]: <Axes: >
```



## RandomForestclassifier

```
In [199]: np.random.seed(42)
          clf = RandomForestClassifier()
          clf.fit(x_train ,y_train)
          clf.score(x_test ,y_test)
```

Out[199]: 0.8695652173913043

```
In [200]: ## y_pred
          y_pred = clf.predict(x_test)
          print(accuracy_score(y_test ,y_pred))
```


0.8695652173913043

```
In [201]: f = confusion_matrix(y_test ,y_pred)
          f
```

Out[201]: array([[ 98, 18],
 [ 18, 142]], dtype=int64)

```
In [202]: sns.heatmap(f ,annot = True)
```

```
Out[202]: <Axes: >
```



## Comparing the scores

```
In [203]: print("svc accuracy score",svc.score(x_test ,y_test))  
print("decision accuracy score",dec.score(x_test ,y_test))  
print("RandomForest accuracy score",clf.score(x_test ,y_test))
```

```
svc accuracy score 0.855072463768116  
decision accuracy score 0.7463768115942029  
RandomForest accuracy score 0.8695652173913043
```

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

## Task 2

In [204]: *## Read the dataset*

```
data = pd.read_csv("../DataSets/Churn_Modelling.csv")  
data.head(5)
```

Out[204]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	

```
In [205]: ##split the data into x and y
df = data.drop(["RowNumber","CustomerId"] ,axis =1)
x = df.drop("Exited",axis = 1)
y = df["Exited"]
x.head(2)
```

Out[205]:

	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58

```
In [206]: y
```

```
Out[206]: 0      1
1      0
2      1
3      0
4      0
..
9995   0
9996   0
9997   1
9998   1
9999   0
Name: Exited, Length: 10000, dtype: int64
```



```
In [207]: ## Now do the preprocessing ( Label Encoding & StandardScaling)

scaling = StandardScaler()
encoder = OneHotEncoder()

category = ["Surname", "Geography", "Gender", "NumOfProducts", "HasCrCard", "IsActiveMember"]
numerical = ["CreditScore", "Age", "Tenure", "Balance", "EstimatedSalary"]

transform = ColumnTransformer([("numerical", scaling, numerical),
                               ("category", encoder, category)], remainder="passthrough")
trans_x = transform.fit_transform(x)
trans_x.shape
```

```
Out[207]: (10000, 2950)
```

```
In [208]: ## Split the data into train test
np.random.seed(42)
x_train ,x_test ,y_train ,y_test = train_test_split(trans_x ,y ,test_size = 0.3 ,random_state = 43)
```

```
In [209]: ## Import necessary libraries

import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense
```

## For relu

In [210]: *## Create a ANN*

```
model = Sequential()

model.add(Dense(32, activation='relu', input_shape=(2950,)))

model.add(Dense(12, activation='relu'))

model.add(Dense(2, activation='softmax'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

model.summary()
```

Model: "sequential\_15"

Layer (type)	Output Shape	Param #
dense_43 (Dense)	(None, 32)	94432
dense_44 (Dense)	(None, 12)	396
dense_45 (Dense)	(None, 2)	26

=====  
 Total params: 94854 (370.52 KB)  
 Trainable params: 94854 (370.52 KB)  
 Non-trainable params: 0 (0.00 Byte)

In [211]: *## Encode y\_labels using onehotencoder*

```
from keras.utils import to_categorical
```

```
y_train_encoded = to_categorical(y_train)
```

```
y_test_encoded = to_categorical(y_test)
```

```
history = model.fit(x_train ,y_train_encoded ,batch_size = 32 ,epochs = 10 )
```

Epoch 1/10

219/219 [=====] - 2s 4ms/step - loss: 0.4698 - accuracy: 0.7993

Epoch 2/10

219/219 [=====] - 1s 4ms/step - loss: 0.3599 - accuracy: 0.8499

Epoch 3/10

219/219 [=====] - 1s 4ms/step - loss: 0.3232 - accuracy: 0.8719

Epoch 4/10

219/219 [=====] - 1s 4ms/step - loss: 0.2902 - accuracy: 0.8876

Epoch 5/10

219/219 [=====] - 1s 4ms/step - loss: 0.2599 - accuracy: 0.8979

Epoch 6/10

219/219 [=====] - 1s 4ms/step - loss: 0.2312 - accuracy: 0.9100

Epoch 7/10

219/219 [=====] - 1s 4ms/step - loss: 0.2100 - accuracy: 0.9186

Epoch 8/10

219/219 [=====] - 1s 4ms/step - loss: 0.1921 - accuracy: 0.9270

Epoch 9/10

219/219 [=====] - 1s 4ms/step - loss: 0.1760 - accuracy: 0.9321

Epoch 10/10

219/219 [=====] - 1s 5ms/step - loss: 0.1624 - accuracy: 0.9367

In [212]: `y_predict_encoded = model.predict(x_test)`  
`y_preds = np.argmax(y_predict_encoded,axis = 1)`  
`model.evaluate(x_test , y_test_encoded)`

94/94 [=====] - 0s 3ms/step

94/94 [=====] - 1s 4ms/step - loss: 0.4874 - accuracy: 0.8093

Out[212]: [0.48743578791618347, 0.809333324432373]


```
In [213]: accuracy_score(y_test ,y_preds)
```

```
Out[213]: 0.8093333333333333
```


```
In [214]: an = confusion_matrix(y_test ,y_preds)
print(an)
sns.heatmap(an ,annot = True)
```

```
[[2138  249]
 [ 323  290]]
```

```
Out[214]: <Axes: >
```



```
In [215]: plt.plot(history.epoch , history.history["loss"])
plt.plot(history.epoch , history.history["accuracy"])
plt.title("Model Accuracy and Loss", fontsize=18)
plt.xlabel("Epoch", fontsize=12)
plt.ylabel("Value", fontsize=12)
plt.legend(["loss" , "accuracy"])
plt.show()
```



## For Tanh

```
In [216]: ## Create a ANN for relu activation

model = Sequential()

model.add(Dense(32, activation='tanh', input_shape=(2950,)))

model.add(Dense(12, activation='tanh'))

model.add(Dense(2, activation='softmax'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

model.summary()
```

Model: "sequential\_16"

Layer (type)	Output Shape	Param #
=====		
dense_46 (Dense)	(None, 32)	94432
dense_47 (Dense)	(None, 12)	396
dense_48 (Dense)	(None, 2)	26
=====		
Total params: 94854 (370.52 KB)		
Trainable params: 94854 (370.52 KB)		
Non-trainable params: 0 (0.00 Byte)		

In [217]: *## Encode y\_labels using onehotencoder*

```
from keras.utils import to_categorical
```

```
y_train_encoded = to_categorical(y_train)
```

```
y_test_encoded = to_categorical(y_test)
```

```
history = model.fit(x_train ,y_train_encoded ,batch_size = 32 ,epochs = 10 )
```

Epoch 1/10

219/219 [=====] - 2s 3ms/step - loss: 0.4502 - accuracy: 0.8053

Epoch 2/10

219/219 [=====] - 1s 4ms/step - loss: 0.3615 - accuracy: 0.8493

Epoch 3/10

219/219 [=====] - 1s 3ms/step - loss: 0.3346 - accuracy: 0.8631

Epoch 4/10

219/219 [=====] - 1s 4ms/step - loss: 0.3035 - accuracy: 0.8769

Epoch 5/10

219/219 [=====] - 1s 4ms/step - loss: 0.2838 - accuracy: 0.8861

Epoch 6/10

219/219 [=====] - 1s 4ms/step - loss: 0.2714 - accuracy: 0.8894

Epoch 7/10

219/219 [=====] - 1s 4ms/step - loss: 0.2628 - accuracy: 0.8873

Epoch 8/10

219/219 [=====] - 1s 5ms/step - loss: 0.2551 - accuracy: 0.8940

Epoch 9/10

219/219 [=====] - 1s 4ms/step - loss: 0.2490 - accuracy: 0.8951

Epoch 10/10

219/219 [=====] - 1s 4ms/step - loss: 0.2476 - accuracy: 0.8973

In [218]: 

```
y_predict_encoded = model.predict(x_test)
y_preds = np.argmax(y_predict_encoded,axis = 1)
model.evaluate(x_test , y_test_encoded)
```

94/94 [=====] - 0s 2ms/step

94/94 [=====] - 0s 2ms/step - loss: 0.4829 - accuracy: 0.8183

Out[218]: [0.48287254571914673, 0.8183333277702332]


```
In [219]: accuracy_score(y_test ,y_preds)
```

```
Out[219]: 0.8183333333333334
```

```
In [220]: an = confusion_matrix(y_test ,y_preds)
print(an)
sns.heatmap(an ,annot = True)
```

```
[[2216  171]
 [ 374  239]]
```

```
Out[220]: <Axes: >
```





In [ ]:

In [ ]: